




Accessibility Benchmark

Aquavena for the Visually Impaired

by **Adrien SALES**



 github.com/adriens/aquavena  linkedin.com/in/adrien-sales

June 14, 2026





Accessibility Benchmark Aquavena for the Visually Impaired

Comparing `adriens.github.io/aquavena` vs `aquavena.nc`
Web Accessibility · Low Vision · WCAG 2.1 AA · New Caledonia

Adrien SALES

 x.com/rastadidi  linkedin.com/in/adrien-sales
 github.com/adriens/aquavena

 5c097f67c44fbc6792b8d81f47e56da4333e7146

Abstract

*This report automatically compares the web accessibility of **adriens.github.io/aquavena** — designed for visually impaired users — against the source website **aquavena.nc**, a dietetic meal delivery service in New Caledonia. Standard CLI tools (*pa11y-ci*, *Lighthouse*, *axe-core*, *webhint*) are used as a common reference. Results show a **100/100 Lighthouse score and 0 WCAG 2.1 AA errors** for our site, versus 91/100 and 67 errors for the official website. All audit data is persisted in a DuckDB database for longitudinal tracking.*

civitech	tech for good	inclusive design	digital inclusion	accessibility	WCAG
visually impaired	social impact	assistive technology	elderly care		
aging population	senior accessibility	human-centered design	design thinking		
New Caledonia					

Table of contents

1	♥ Inception	6
2	📊 Executive Dashboard — Delivery KPIs	6
2.1	EVM — Planned Value vs Earned Value	6
2.2	Strategic impact — multi-axis radar	8
2.3	Category delta — gap closed per accessibility dimension	8
2.4	Delivery velocity — engineering throughput per release	9
2.5	DORA metrics — elite performance	10
2.6	Development intensity	11
2.7	AI-assisted delivery — Claude token consumption	11
2.8	Semantic analysis of commit messages	13
2.9	AI-leverage ratio per cluster — where the AI actually pays off	25
3	🎯 Goals and Broader Vision	26
3.1	Immediate goal	26
3.2	A replicable methodology	26
3.3	Transferability to other accessibility domains	27
3.4	From measurement to optimisation	29
3.5	Accessibility score as Business Value in Scrum	29
4	🔗 Site Architecture	30
4.1	Design principles	30
4.2	Technology stack	30
4.3	Data pipeline	31
4.4	Accessibility engineering	31
5	🔧 Tools	32
6	🔍 Methodology	33
6.1	Pipeline overview	33
6.2	Release discipline	34
6.3	Tool coverage	35
6.4	Tools	36
6.5	Standard	36
7	📊 Results — Our Site	36
7.1	pa11y-ci — WCAG 2.1 AA	36
7.2	Lighthouse Accessibility	37
8	🌐 Results — aquavena.nc	38
9	⚖️ Comparison	38
9.1	HTML quality & page weight	38
9.2	Privacy & user tracking	39
10	🛡️ Accessibility Features	40

11 🏠 Lighthouse History	41
12 📊 Score by Version	41
13 🗄 Database	42
13.1 Schema	43
13.2 Example queries	45
14 📖 Glossary	45
15 🏠 WebMCP — The Site as an Agentic Interface	47
15.1 Why WebMCP?	47
15.2 Implementation	47
15.3 What It Brings	48
16 🗺 Future Work — ML & Graph Data Science Roadmap	48
16.1 Scale to N sites	48
16.2 Graph data science on the audit network	48
16.3 Deepen the commit-semantics work	49
16.4 Time-series of accessibility evolution	49
16.5 AI-Agent Universal Accessibility Pipeline	49
16.6 Honest data requirements	50
16.7 New perspectives unlocked by the 146-commit dataset	51
17 📖 Further Reading	52
18 🚩 Conclusion	52
A 📖 Appendix — Full Table Descriptions	53
A.1 git_commits	54
A.2 ci_runs	54
A.3 pally_issues	55
A.4 lighthouse_audits	56
A.5 score_history	56
A.6 gh_releases	57
A.7 html_quality	58
A.8 privacy_audit	58
A.9 claude_sessions	59
A.10 claude_usage	60
A.11 model_pricing	60

1 ♥ Inception

This project was not born in a research lab or a corporate accessibility audit. It started at a kitchen table, from a very simple and concrete situation.

My mother has been a loyal Aquavena customer for years. Every week she orders her meals from their service, and she genuinely enjoys the food. But she is visually impaired — and as her condition evolved, consulting the weekly menus on the official website became increasingly difficult, then impossible to do on her own. The layout, the font sizes, the colour contrasts, the absence of any audio alternative: none of it was designed with her in mind.

The consequence was small but recurring, and that is precisely what made it significant. Every week, she had to wait for someone to be available — a family member, a neighbour — to read the menus out loud to her. She lost the freedom to plan her own meals at her own pace. A minor dependency, perhaps, but one that repeated itself week after week, quietly eroding her autonomy.

I am a software engineer. I realised that a few hours of work could give that autonomy back to her: a dedicated page she could open on her tablet, with large text, high contrast, a dark mode for tired eyes, and a button to have the menus read aloud in her own language. Something she could use alone, at any hour, without asking anyone for help.

That is why this site exists. It is not a product, not a startup, not a statement. It is a small act of care, built with the tools I know, for the person who matters most. The accessibility work documented in this report is a direct consequence of that intent: if the site is meant to help someone with low vision, it had better actually be accessible.

A note to the Aquavena team, should they ever read this: there is no competition here, no monetisation, no hidden agenda. Every menu on this site links back to aquavena.nc to place orders. This is simply a more accessible front door to a service my mother loves.

2 📊 Executive Dashboard — Delivery KPIs

This dashboard reads directly from `benchmark.duckdb` at render time and produces the KPIs most relevant to a delivery committee: classical **EVM** (Earned Value Management), modern **DORA** (DevOps Research and Assessment) elite-performance indicators, the **WCAG defect burndown**, the **sprint velocity** profile, and the **strategic gap** between our work and the source site. Nothing in the charts below is hardcoded — every figure is computed live from the six FK-linked tables described in the appendix. Re-running the pipeline regenerates the dashboard automatically.

2.1 EVM — Planned Value vs Earned Value

A linear ramp from the **aquavena.nc baseline (91)** to the **target score (100)** defines the *Planned Value* curve — the score the project would have had at each release if it had improved at a constant rate. The *Earned Value* curve is the **actual Lighthouse score**


recorded at each release tag. Their ratio is the **Schedule Performance Index (SPI = EV / PV)**, the most-watched indicator in PMI's Earned Value Management framework.

2.1.1 How to read the SPI chart

The dashed grey line at **SPI = 1.0** is the *plan baseline*. Every bar grows **upward from that line** to the actual SPI value for that release:

- **Teal bar above 1.0** → release is **ahead of plan** (Earned Value > Planned Value). The taller the bar, the more headroom recovered relative to what was scheduled.
- **Red bar below 1.0** (none in this dataset) → release would be **behind plan**. The longer the bar, the further behind.
- **No visible bar** (exactly at 1.0) → release is **on plan**.


For this project the chart shows a **uniformly-teal “ahead of plan” pattern**: the very first release (v0.1.1) lands at **SPI ≈ 1.085** — already 8.5% ahead of where the linear plan said we should be. The lead narrows release after release as the plan baseline catches up to the actual score, ending exactly on target at v1.7.0 (SPI = 1.000). This is the visual signature of a project that **set the bar at the maximum value from day one** and held it across every checkpoint — the cleanest possible EVM trace for a PMI committee.



2.2 Strategic impact — multi-axis radar


Each Lighthouse audit is classified into one of five accessibility categories, and the **pass rate per category** is computed per site. The radar makes the strategic gap immediately readable: where the source site falls short, where it does well, and where our rebuild closes the gap entirely.

Strategic impact ... pass rate by accessibility category
% of Lighthouse audits passed per category, per site




2.3 Category delta — gap closed per accessibility dimension

The replay shows our site at **100/100 Lighthouse, 0 pa11y errors** across every tag, while aquavena.nc sits at 91/67. Instead of an iterative burndown, the chart below visualises the **one-shot transformation** as a category-by-category **gap closed**: for each accessibility dimension Lighthouse audits, the bar shows the pass-rate delta between the source site and ours. The longer the bar, the bigger the gap the rebuild closed in that dimension.



2.4 Delivery velocity — engineering throughput per release


Per-release **lines of code changed** (teal bars) — the engineering throughput each sprint poured into the codebase. The **cumulative lines delivered** since v0.1.1 is overlaid as an amber line on the right-hand axis. Reading the chart: after the initial scaffold (v0.1.1), the project sustained a steady cadence of small surgical sprints early on (v0.1.2–v0.1.5), then larger investment releases for tooling and reporting (v1.5.0–v1.7.0).



2.5 DORA metrics — elite performance

The **DORA** acronym used throughout this report stands for **DevOps Research and Assessment** — the Google research programme that defines the four key software-delivery metrics (Deployment Frequency, Lead Time for Changes, Change Failure Rate, MTTR). This is distinct from the EU financial regulation of the same acronym (Digital Operational Resilience Act, 2022/2554).


The DORA metrics (DevOps Research and Assessment) are computed from the timestamps stored in `gh_releases.published_at_ts` and from the `ci_runs` table. Results land in the **Elite** tier as defined by Forsgren et al., *Accelerate State of DevOps Report 2023* (Google Cloud / DORA, dora.dev/research/2023).



Forsgren et al., *Accelerate State of DevOps Report 2023*, Google Cloud / DORA — Elite tier: on-demand deploys, lead time < 1 h, MTTR < 1 h.

Lead Time for Changes — time between consecutive releases (commit merged → live on GitHub Pages). Measured here as the median gap between release timestamps. A short lead time means fixes and features reach users quickly.

2.6 Development intensity



2.7 AI-assisted delivery — Claude token consumption

This project was built with **Claude Code**. Every assistant message and every cached context replay is recorded locally; `extract_claude_usage.py` parses those transcripts into the `claude_sessions` and `claude_usage` tables. Joining with `git_commits` via timestamp range gives **tokens per release** (apportioned by the time elapsed between consecutive tags).

i ⓘ What does “cache hit” mean here?

Claude charges for **every** token, but the **prompt cache** changes the rates dramatically:

Token category	Price (Sonnet 4.6)	Price (Opus 4.7)	Relative
Fresh input (cache miss)	\$3.00 / M	\$15.00 / M	1.0×
Cache write (one-off setup)	\$3.75 / M	\$18.75 / M	1.25×
Cache read (cache hit)	\$0.30 / M	\$1.50 / M	0.1×
Output tokens	\$15.00 / M	\$75.00 / M	5.0×

A **cache hit** is when the model serves part of the prompt from cache instead of re-encoding it. It is **still billed**, but at **10% of the fresh-input rate**.

Why our 98% cache-hit ratio matters. In a long Claude Code conversation the entire prior transcript is replayed as context with every new message. Without caching, this re-encoding cost would scale roughly quadratically with session length. With caching, it scales linearly — and at the 10% rate. Concretely: the ~420 M cache-read tokens we logged would have cost **\$1,260** at fresh-input rates; they actually cost about **\$126**. A **10× saving**, automatically.

What “good” looks like in our case. A high cache-hit ratio ($\geq 90\%$) means we are working as efficiently as possible — long, coherent sessions reuse the same context. A **low** cache-hit ratio would indicate either many short sessions (each pays the cache-write tax without amortising it) or the conversation constantly bringing in fresh files. For project-style work like this benchmark, the 98% ratio is **the ceiling, not the average** — it is the strongest signal that the project was delivered in a small number of focused sessions rather than fragmented across many cold starts.



272.67 M

Sonnet 4.6
tokens



404.97 M

Opus 4.7
tokens




97.8%

Cache hit
ratio



2497 | 3789



Messages
you | Claude



2.8 Semantic analysis of commit messages

All **146 commit messages** were embedded with the BAAI/bge-m3 model (1024-dimensional **multilingual** embeddings — handles the project’s mix of French and English commit messages), clustered with K-Means into 6 groups, and projected to 2D with PCA. The result is stored in two tables — `commit_classification` (per-commit type/scope/cluster/coords) and `commit_clusters` (per-cluster size and characteristic terms) — and is queryable directly from DuckDB.

The chart on the left shows the **conventional-commit type mix** (parsed from the `feat: / fix: / docs: /` prefixes); the chart on the right is the **2D semantic map** — each dot is one commit, coloured by its semantic cluster.



The cluster terms below come from per-cluster TF-IDF; each row's exemplar is the commit closest to the cluster centroid in embedding space.

TF-IDF (Term Frequency - Inverse Document Frequency) measures how characteristic a word is to one cluster compared to all others. A term scores high when it appears often *inside* a cluster (TF) but rarely across the other clusters (IDF). It is what makes “accessibility” a top term for the feature cluster but not for the fix or chore clusters — even if the word appears everywhere, the IDF weight penalises it unless it is truly distinctive.

Cluster	Commits	Top terms	Exemplar
0	26	feat, add, feat add, about, rss	feat: add RSS and iCal sections to About page
1	24	style, pour, html, de, header	style: agrandir la taille du texte pour meilleure lisibilité
2	14	benchmark, feat benchmark, feat, section, ml	feat(benchmark): section 'Site Architecture' (stack, principes, accessibilité)
3	8	aquavena, aquavena-sdk, de contact, contact, feat	feat: infos de contact Aquavena (adresse, tél, FB, Instagram, Play Store)
4	8	fix, contraste, fix contraste, wcag, npm	fix(a11y): corrige toutes violations contraste restantes + structure benchmark/
5	3	section valentin, valentin, lien, dans, section	feat(about): lien site AVH comité Nouvelle-Calédonie dans section Valentin Haüy

2.8.1 Graph data science — anchors & bridges

The embeddings already cluster the commits. Treating them as a **graph** — nodes = commits, edges = cosine similarity ≥ 0.55 — adds a second layer of insight: which commit best *represents* each theme (highest **PageRank** within its cluster, the **anchor**), and which commits *connect* different themes (highest **betweenness centrality**, the **bridges**). Both metrics are computed with `networkx`, persisted in `commit_graph_metrics` for direct SQL queries, and exported as `benchmark/data/commit_graph.gexf` — a **Gephi-ready file** with all node attributes (cluster, PageRank, betweenness, degree, type, scope, timestamp) and edge weights, so the graph can be explored interactively (ForceAtlas2 layout, community detection, etc.) in Gephi without re-running the pipeline.




Figure 1: Commit-similarity graph — **ForceAtlas2** layout (Gephi native algorithm, fa2-modified port, 1 400 iterations) · node colour = semantic cluster · node size = PageRank (percentile-scaled, $\approx 80\times$ area ratio) · edges = cosine similarity ≥ 0.55




Figure 2: Cluster-split view — same ForceAtlas2 positions as above · each panel spotlights one semantic cluster (full glow) while the rest fade to grey · intra-cluster edges are coloured, bridge edges are muted · anchor label = commit closest to centroid

Table 1: Anchor commit per cluster — the most-representative work in each theme

Cluster	PageRank	Anchor commit
C0	0.0296	feat: move accessibility score section to top of About page
C1	0.0253	style: agrandir la taille du texte pour meilleure lisibilité
C2	0.0261	feat(benchmark): section 'Site Architecture' (stack, principes, accessibilité)
C3	0.0130	chore: update repo URLs after GitHub rename to aquavena-sdk
C4	0.0121	fix(a11y): contraste onglets régimes en mode nuit
C5	0.0107	feat: bouton menu du jour, lien À propos dans header, section Valentin Haüy

Table 2: Top-5 bridge commits — connect distinct themes (highest betweenness centrality)

From cluster	Betweenness	Bridge commit
C0	0.1140	feat: add A+/A- and dark mode to About page header
C1	0.0972	style: agrandir la taille du texte pour meilleure lisibilité
C0	0.0891	feat: add About page explaining the accessibility approach
C0	0.0884	feat: move accessibility score section to top of About page

C1	0.0775	feat(report): refonte du diagramme Universal Accessibility Pipeline
----	--------	---

i What the graph reveals

Anchors (cluster representatives). Across all six clusters, the anchor is almost always an **About-page feature commit** (feat: add A+/A- and dark mode to About page header, feat: move accessibility score section to top of About page, feat: add read-aloud button...). This is the structural finding: the **About page is the semantic centre of the codebase** — the single page where the project explains itself most fully, and where the embeddings find the highest density of similar content.

Bridges (theme-linkers). The top-5 betweenness commits all touch **multiple themes simultaneously** — they're the commits that introduced A+/A- and dark mode (style + feature), or moved the accessibility-score section to the About page (a11y + About). These are the **integration commits** — the ones a knowledge-graph would surface as architecturally critical.

Why this matters at scale. With 82 commits the graph is small, but the *technique scales linearly with commit count*. Applied to a 1000-commit codebase, PageRank surfaces the commits that anchor each subsystem; betweenness surfaces the integration commits whose removal would *fracture* the codebase. Both are signals an experienced reviewer would notice manually — this method makes them queryable.

2.8.2 Cluster drift — a project maturity signal

What it is. Each commit belongs to one of six semantic clusters discovered earlier. The **cluster drift** chart plots, for every release tag along the project's history, the *proportion* of commits from each cluster that landed in that release. Stacked together, these proportions form a **time-series of the project's semantic mix**.


What it measures. The shifting composition of work, release after release. Early releases tend to be dominated by **core-feature clusters** (accessibility fixes, UX features); mid-life releases shift toward **hardening clusters** (dependency bumps, bug fixes, style); mature releases drift toward **meta-work clusters** (reporting, documentation, tooling, agentic instrumentation). This is the **maturity signature** of a software project.

What it is for. A *zero-label* maturity metric. No human annotation is required: the cluster mix per release falls out of conventional commits and sentence-transformer embeddings already computed by the pipeline. It answers questions a PMI committee or steering board typically asks subjectively — “*is this project still in feature mode, or is it in maintenance? has the team moved on to a new phase?*” — and turns them into a measurable, reproducible signal.

What to watch for.

- A **steady single colour** across releases → monolithic project with one type of work being done over and over.

- A **smooth gradient** from one cluster to another → healthy evolution from build to stabilisation.
- An **abrupt cluster appearing late** → the project pivoted (new domain, new platform, new audience). On Aquavena, the *Benchmark & ML report* cluster appears massively from v1.7.0 onwards: this is the visual signature of the project shifting from *building an accessible site* to *measuring and publishing what was built* — the report you are reading right now.



Reading this chart on Aquavena.

- **v0.1.2-v0.1.4** (the very first releases) are 100% *WCAG fixes* — the foundational accessibility work: contrast corrections, pa11y audit setup, WCAG 2.1 AA compliance from the ground up. This is the *raison d'être* of the project, visible in its purest form.
- **v0.1.5** is the **first multi-cluster release**: WCAG work is still present, but *Style & dark mode*, *Site features & About* and *Benchmark & ML report* clusters all appear simultaneously. This is the moment the project diversified — from one concern (make it WCAG-clean) to four (style polish, feature additions, benchmark scaffolding).
- **v1.7.0** stands out as the only mid-life release where the *Site features & About* cluster occupies a large share — this matches the historical record (RSS, iCal, About-page content all landed around this tag).
- **v1.9.0 → v1.13.x** are dominated by the *Benchmark & ML report* cluster — the long **report-engineering phase**: this is where the 100-page PDF you are reading was built. Feature work shifted to a data-science/reporting register.
- **v1.12.1** is the only release of this phase where *WCAG fixes* re-emerge at parity with the report work — likely a regression catch that triggered contrast revalidation.

- **v1.14.0** closes the chart with a return to 100% WCAG *fixes* — the **Tailwind v4 migration** release: a platform overhaul whose primary risk was contrast regression, so most of the work was re-validating accessibility.


Note: the chart stops at v1.14.0 because newer release tags (v1.15.0–v1.24.0) — including the entire WebMCP & agentic-pipeline work documented in this report — were published *after* the gh_releases snapshot used by this build. Re-running the pipeline will extend the timeline and reveal the *Site features & About* and *Benchmark & ML report* clusters dominating the rightmost releases — the agentic pivot’s signature.

This single chart tells the story of the project’s life cycle without a single line of manual annotation. Applied to other codebases, it becomes a portable **maturity radar** — a public-service site stuck on a single cluster for two years is not evolving; a site whose cluster mix shifts is iterating.

2.8.3 BUILD vs RUN — hackathon intensity over time

The chart below answers three questions at once: **when** was effort delivered, **what kind** of work it was (BUILD = new value, RUN = maintenance), and **how intense** the sessions were.

Why conventional commits with scope make this possible. A plain fix: or docs: tells you the type but not the topic. The scope field is what makes classification precise and automatic: docs(benchmark): lands in BUILD because the benchmark report is a deliverable; a bare docs: (release notes, README) stays in RUN. **No time-tracking tool was filled in. No spreadsheet was maintained.** The metadata already lived in every commit message — the pipeline simply reads it. This is the core argument for conventional commits: they turn git history into a queryable KPI source with zero developer overhead.



Reading the chart. Each bar is the sum of time-spent (inter-commit gaps capped at 30 min) falling in a 2-hour window. A tall bar means a focused session; a flat stretch means a break or a day off. The rug marks at the bottom of each lane show individual commits coloured by type — the mix of colours reveals whether a burst was feature-driven (feat, teal) or fire-fighting (fix, red).

The queryable source of truth. The active-time estimate underlying the bars is directly queryable:

```
WITH ordered AS (
  SELECT CASE WHEN commit_type = 'feat' THEN 'BUILD'
              WHEN commit_type = 'docs' AND scope = 'benchmark' THEN 'BUILD'
              ELSE 'RUN' END AS activity,
         gc.commit_ts,
         LAG(gc.commit_ts) OVER (ORDER BY gc.commit_ts) AS prev_ts
  FROM commit_classification cc JOIN git_commits gc USING (git_sha)
)
SELECT activity,
       ROUND(SUM(LEAST(1800,
                       GREATEST(0, EXTRACT(EPOCH FROM (commit_ts - prev_ts)))
                       )) / 60.0, 1) AS minutes_active
FROM ordered GROUP BY activity ORDER BY activity;
```

i ≡ What the BUILD/RUN split reveals


Conventional commits eliminate time-tracking overhead. The BUILD/RUN classification is derived entirely from `commit_type` and `scope` — two fields that every developer writes anyway. No Jira worklog. No Toggl timer. No spreadsheet. The git history is the timesheet, and the pipeline reads it automatically. This is the strongest practical argument for adopting conventional commits on any project: they turn a developer's natural commit habit into a queryable KPI source at zero extra cost.

The scope field is the precision layer. A bare `docs(benchmark):` commit would be misclassified as RUN. `docs(benchmark):` correctly lands in BUILD because the scope carries the topic. Without scope, the classification degrades to a coarse type-only split. With scope, it is surgically accurate — and entirely machine-readable.

The marathon is visible. The intensity bars do not lie: the hackathon weekend appears as a sustained sequence of tall BUILD bars, not a polite 9-to-5 profile. That kind of effort is invisible in a story-point count or a commit-count chart; it surfaces only when you measure *time actually spent at the keyboard*.


2.8.4 Tokens consumed by category

Joining `commit_classification` with the `commit_tokens` view gives the **AI delivery cost per category** — total tokens spent on every `feat:`, every `fix:`, every semantic cluster, etc. This is the **first ML-adjacent insight from the real data**: it reveals which kinds of work are token-cheap versus token-expensive.



By commit scope ... where the report-writing effort lives

Scopes from feat(scope): / fix(scope): ... 'benchmark' = analysis report · '(none)' = no scope declared




325 M

Writing this report

*tokens spent on
scope=benchmark
(32 commits)*


**5.22 M >
5.90 M**

Fix avg > feat avg
*+12% more tokens per
fix*


2.12 M

Style cluster
avg / commit (C0)


6

Clusters found

*unsupervised,
k-means*

i 🔑 Key findings — AI delivery cost vs work category

1. The benchmark/reporting strand dominates AI usage. Cluster C1 (benchmark, feat benchmark, section) consumed **163 M tokens across 17 commits — an average of 9.6 M tokens per commit**, three to six times the per-commit average of any other cluster. The same picture appears under the **supervised lens**: commits explicitly tagged scope=benchmark (12 commits) consumed **144.8 M tokens at an average of 12.07 M per commit** — by far the biggest scope in the project. Whichever filter you apply, the measurement and reporting infrastructure (this report, the audit scripts, the FK schema, the dashboard) is *the single most AI-intensive strand of the project*. The user-facing site itself was comparatively cheap to produce.



2. Fixes cost more tokens per commit than features. fix: commits average **4.72 M tokens each** versus feat: at **3.82 M** — **24% more deliberation per fix** despite fewer commits (13 vs 42). This matches engineering intuition: a fix requires diagnosing a defect *before* addressing it, while a feature can be specified and implemented in one shot. For accessibility work specifically, fixes also require validating against multiple audit tools — pa11y, Lighthouse, the W3C HTML validator — which multiplies the verification cost.

3. Style commits are the cheapest category. Cluster C0 (style, feat, la, du) averages only **1.59 M tokens per commit** — the lowest per-commit cost in the project. UI/style tweaks are short feedback loops with immediate visual confirmation, so they need fewer tokens to land. *Implication for AI-assisted teams: deliberately split work into style-grade chunks where possible to minimise cost.*

4. The unsupervised clustering rediscovered the supervised labels. The K-Means clusters built from sentence-transformer embeddings — i.e. with **zero knowledge of the feat:/fix:/docs: prefixes** — recovered the same structural separation. C2 is essentially the fix: cluster, C1 the feat(benchmark): cluster. *This is a validation that the embeddings encode semantic meaning faithfully*; it also means the technique would still work on commit history without conventional-commit discipline.

5. The flat 1.66 M baseline is an apportionment artifact, not a finding. docs, chore, and ci all show the **same** average per commit (1.66 M). This is because those commits cluster very tightly in time and the apportionment view distributes session tokens evenly across same-second commits. *With a larger dataset (more sessions, more spread-out commits) this artifact disappears.* Disclosed here for transparency.

2.9 AI-leverage ratio per cluster — where the AI actually pays off



i 📌 What the AI-leverage view reveals

A 5.5× gap in AI-intensity between the most and least token-heavy clusters. The *WCAG fixes* cluster consumes **11.7 M tokens per commit** on average, while *Site features & About* consumes only **2.1 M tokens per commit**. This is not noise — the gap reflects a real **cognitive depth difference** between the work types: some commits are written in one shot, others require sustained iteration with the assistant.

The Opus dependency ratio is a “premium model” signal. *WCAG fixes* routed **73%** of its tokens to Opus 4.7 — the project’s most expensive model — while *Site features & About* used Opus only **28%** of the time. Clusters with a high Opus share are the ones where the developer felt that Sonnet’s output quality was *not enough* and consciously escalated to Opus. **It’s a direct measurement of which problem domains are intrinsically hard, derived purely from model-routing telemetry.**

Where the IA budget actually went. The *Benchmark & ML report* cluster represents the largest absolute IA spend — **\$441** across 33 commits. This number is a **directly actionable PMI signal**: a delivery committee can now see, on a per-cluster basis, where one cluster cost $N\times$ more than expected — a maturity proxy, a complexity proxy, and a budget proxy in a single number.

Why this matters beyond Aquavena. No team-wide instrumentation, no JIRA fields, no developer self-reporting was used. The numbers above are entirely derived from `commit_tokens` and `commit_classification`, both populated automatically by the pipeline. Replicated on any AI-assisted codebase, this view gives a **per-cluster ROI on AI spending** — the missing piece between “we spent \$X on Claude this month” and “we spent it on Y type of work that produced Z value”.

3 🕒 Goals and Broader Vision

3.1 Immediate goal

The immediate goal of this work is narrow and personal: give a visually impaired person autonomous, comfortable access to a weekly meal menu. That problem is solved. But the process of solving it surfaced something more general.

3.2 A replicable methodology

Building this site required answering a question that turns out to be surprisingly hard in practice: *how do you know whether a website is actually accessible?* Subjective assessment does not scale. Manual inspection is inconsistent. What this project demonstrated is that a fully automated, score-based pipeline can answer that question reproducibly and cheaply.


The methodology is as follows. A set of open-source CLI tools — `pa11y-ci`, `Lighthouse`, `axe-core` — each consume a URL and produce a structured output: a JSON file containing a score, a list of violations, and enough metadata to track change over time. Those outputs

are loaded into an analytical database (DuckDB), which makes them queryable, comparable, and archivable. A report is then generated automatically from that database.

The key insight is that **accessibility becomes measurable**. Once it is measurable, it can be benchmarked, monitored, and improved in the same way that software performance or test coverage is managed: with targets, regressions, and trends.


The automated accessibility improvement cycle

PDCA (Plan-Do-Check-Act) wheel · driven by GitHub CI/CD



3.3 Transferability to other accessibility domains

Low vision is one accessibility challenge among many. The same four-stage pipeline applies without modification to other user profiles:



- **Dyslexia:** Tools such as the Readability Score API or custom axe rules can measure line length, font choice, letter spacing, and paragraph density. A dyslexic-friendly site has computable characteristics, and a score can be assigned to them just as Lighthouse assigns a score to colour contrast.
- **Motor impairment:** Keyboard navigability, focus order, and touch target size are already audited by Lighthouse and axe-core. A pipeline targeting these specific rules would produce a “motor accessibility score” for any site.
- **Cognitive load:** Metrics such as reading level (Flesch-Kincaid), page complexity, and number of interactive elements per viewport can be computed automatically and tracked over time.
- **Attention / focus difficulties:** Profiles affected by attention disorders benefit from a different set of computable signals — text density per viewport, animation count, auto-playing media, visual noise (number of competing focal points), and reading-progress indicators. The signals exist; they just need to be wired into the same pipeline. This need became immediately visible to me while teaching such a student: the same course material that worked for everyone else was actively counter-productive for them, and small adaptations made the difference between disengagement and full participation.
- **Autism spectrum:** People on the autism spectrum often need clear, predictable interfaces with reduced sensory load — minimal motion, explicit rather than implicit interactions, consistent vocabulary, and tightly predictable feedback. Many of these properties are also computable: prefers-reduced-motion honoured, animation duration thresholds, layout stability score (CLS), interaction predictability heuristics. The pipeline applies without modification.

The template is always the same: identify the user profile, map it to computable signals, plug those signals into the pipeline, and let the data speak. The Aquavena project is one instantiation of that template. Others could follow with minimal adaptation — including, potentially, the official aquavena.nc site itself.

3.4 From measurement to optimisation

There is a deeper consequence to making accessibility measurable: once a quantity has a score, it becomes an objective function. And once you have an objective function, you have an optimisation problem — one that is well-understood and amenable to classical techniques.

Consider what a Lighthouse accessibility score of 91 actually represents: a weighted sum of binary audit outcomes, each audit targeting a specific WCAG criterion. Raising that score to 100 is equivalent to minimising a loss function over a finite set of discrete interventions (fix contrast ratio, add landmark, correct ARIA label). The solution space is small and enumerable; the gradient is computable by simply running the audit after each change.

This mirrors problems in other engineering disciplines. Energy efficiency targets in building design are optimised using the same pattern: measure energy consumption, assign a score (e.g. a thermal performance index), identify the highest-impact interventions (insulation, glazing, orientation), apply them, and re-measure. The mathematics are identical; only the domain changes.

The implication for accessibility is significant. Rather than treating WCAG conformance as a compliance checkbox to be reviewed once, it can be treated as a continuous optimisation target — monitored in CI/CD pipelines, tracked over time in a database, and reported automatically. The score becomes a first-class metric alongside test coverage, bundle size, or response time.

3.5 Accessibility score as Business Value in Scrum

The optimisation framing has a direct consequence for agile teams. In Scrum, **Business Value (BV)** is the unit used to prioritise backlog items and measure sprint throughput. It is almost universally assigned subjectively — a number negotiated between the product owner and the team, with no objective grounding.

The accessibility score changes that. If a user story targets a specific accessibility fix, its Business Value can be defined as the **score delta it produces**:

$$BV(\text{story}) = s_{\text{after}} - s_{\text{before}}$$

A story that raises the Lighthouse score from 91 to 96 carries a BV of 5. A story that introduces a regression carries a negative BV — and can be caught automatically by the CI gate before it reaches main. Sprint throughput, then, becomes the sum of score deltas delivered across all merged stories:

$$\text{Throughput}_{\text{sprint}} = \sum_i \Delta s_i$$

This has three practical consequences for a Scrum team:

- **Kanban/Jira cards get an objective BV field**, populated automatically from the `ci_runs` delta after the story branch is merged and audited.
- **Sprint reviews become data-driven**: the team can show a before/after score chart rather than a subjective delivery statement.
- **Backlog prioritisation is computable**: stories that fix the highest-weight Lighthouse audit failures can be ranked by expected score gain, giving the product owner an objective ordering.

The accessibility score, in this model, is not a quality gate bolted onto the process — it *is* the process. It becomes the real unit of value delivery.

4 </> Site Architecture

4.1 Design principles

The companion site was built around a single constraint: every design decision had to serve a visually impaired user first. Three principles guided the work throughout.

Readability before aesthetics. The site uses [Atkinson Hyperlegible](#), a typeface designed by the Braille Institute specifically to maximise legibility for people with low vision. Each letterform is engineered to be unambiguous: no `l` that looks like `1`, no `0` that looks like `o`. The base font size is 18px, adjustable at runtime via `A+` / `A-` buttons that persist across sessions in `localStorage`.

Multiple access modalities. Users should never be forced to read. Every menu can be listened to via the Web Speech API (text-to-speech), one day at a time or the entire week in sequence. Menus are also exposed as RSS feeds and iCal calendars so that a screen reader, a podcast app, or a calendar assistant can consume them without opening a browser at all.

Progressive Web App. The site is installable on any device (iOS, Android, desktop) and works fully offline once cached, so connectivity issues — relevant in parts of New Caledonia — do not block access to the weekly menus.


4.2 Technology stack

Layer	Technology	Role
Framework	Astro 4	Static site generator — zero JS by default
Styling	Tailwind CSS 3	Utility-first CSS, WCAG-compliant colour palette
Typography	Atkinson Hyperlegible	Typeface designed for low vision

Icons	Font Awesome 6	Accessible SVG icons with aria-labels
TTS	Web Speech API	Browser-native text-to-speech, no dependency
PWA	Vite PWA / Workbox	Service worker, offline cache, installability
Data	Python + SDK	Custom scraper — fetches menus from aquavena.nc
Feeds	Astro RSS	RSS 2.0 feeds per diet regime
Calendars	ical.js	iCal (.ics) exports per diet regime
CI/CD	GitHub Actions	Nightly data fetch + build + deploy to GitHub Pages
AI interface	Claude MCP skill	Natural-language menu queries — MCP server hosted on Hugging Fa
Task runner	go-task	Taskfile.yml — one command drives the full build/audit/report pipelin
DB tooling	SchemaCrawler	Reads DuckDB metadata via JDBC and generates FK-aware ER diagra

4.3 Data pipeline

Menus are not hardcoded. A Python SDK scrapes the weekly menus from aquavena.nc each night via a GitHub Actions workflow. The scraped data is written to `src/data/menus.json`, which Astro consumes at build time to generate the static HTML pages. The resulting site is deployed to GitHub Pages with no server-side component.



The site is therefore always up to date by the following morning, with no manual intervention required.

4.4 Accessibility engineering

WCAG 2.1 AA conformance was achieved iteratively, guided by three complementary tools run after every change: **pa11y-ci** (WCAG rule violations), **Lighthouse** (scored audit), and **axe-cli** (axe-core engine). The main categories of issues identified and resolved were:

- **Colour contrast (1.4.3):** All interactive elements — speak buttons, tab labels, action links — were systematically audited. Colours such as bg-orange-500 and bg-sky-500 were replaced with their -700 variants to meet the 4.5:1 ratio for normal text and 3:1 for large bold text.
- **Landmark regions (4.1.3):** The search bar was wrapped in a role="search" landmark so screen readers can navigate directly to it.
- **Label in name (2.5.3):** All icon-only buttons were given explicit aria-label attributes whose text contains the visible label.
- **Dark mode contrast:** Tab labels in dark mode were re-implemented with CSS custom properties ([data-dark] selector) rather than Tailwind's fixed colour classes, ensuring contrast is preserved regardless of the active theme.

5 🦄 Tools

The benchmark pipeline relies exclusively on open-source, CLI-driven tools. Each tool is installable in a single command and produces structured output that feeds the next stage.

Tool	Role	Install
pally-ci	Batch WCAG 2.1 AA audit via headless Chrome; reports errors, warnings, and notices per URL	<code>npm i -g pally-ci</code>
Lighthouse	Google's scored accessibility audit (0-100); covers contrast, ARIA, keyboard nav, structure	<code>npm i -g lighthouse</code>
axe-cli	Deque's axe-core engine via CLI; complements Lighthouse on ARIA semantics	<code>npm i -g axe-cli</code>
webhint	HTTP headers, security, performance, and accessibility hints	<code>npm i -g hint</code>
DuckDB	Portable single-file analytical SQL database; stores all audit results	<code>brew install duckdb</code>
SchemaCrawler	Reads database metadata and generates ER diagrams (used for the schema diagram in this report)	<code>brew install schemacrawler</code>
Graphviz (dot)	Renders SchemaCrawler's DOT output to PNG/PDF	<code>apt install graphviz</code>


Tool	Role	Install
Quarto	Scientific publishing system; executes R code and renders this document to PDF via XeLaTeX	quarto.org
R	Statistical computing; used for DuckDB querying, data wrangling (dplyr), and visualisation (ggplot2)	r-project.org
go-task	Task runner (Taskfile.yml); orchestrates the full pipeline with a single command	<code>brew install go-task</code>

All Node.js tools require **Node 18+** and are pinned via `package.json` in the `site/` directory. R package dependencies are installed via `task setup:r-deps`.

6 Methodology

6.1 Pipeline overview

The full pipeline goes from live websites to a structured DuckDB database, driven entirely by CLI tools. Each step produces a JSON artefact that feeds the next.



6.2 Release discipline

Each accessibility fix is expressed as a **conventional commit** — a structured commit message whose prefix encodes the nature of the change:

Prefix	Meaning	Version bump
fix:	Bug fix or WCAG violation corrected	Patch (1.0.0 → 1.0.1)
feat:	New accessibility feature added	Minor (1.0.0 → 1.1.0)
perf:	Performance improvement	Patch

Prefix	Meaning	Version bump
BREAKING CHANGE:	Structural redesign	Major (1.0.0 → 2.0.0)

On every merge to main, a GitHub Actions workflow parses the commit history, determines the appropriate semantic version bump, tags the repository (v1.2.3), and publishes a **GitHub Release** with an auto-generated changelog listing every fix by its commit message.

This discipline has a direct consequence for the improvement loop: every version tag is a checkpoint. The `git_sha` column in `ci_runs` ties each audit score to an exact commit, and each release to a before/after score delta:

```
-- Score delta between two consecutive releases
SELECT curr.git_sha, curr.lh_score - prev.lh_score AS delta
FROM ci_runs curr
JOIN ci_runs prev ON curr.rowid = prev.rowid + 1
WHERE curr.triggered_by = 'ci';
```


The release is therefore the *unit of measurement* in the optimisation loop — not a deployment formality, but a versioned evidence of improvement.

6.3 Tool coverage

No single tool covers the full accessibility surface. The chart below maps each tool against the dimensions it measures, showing why the combination is necessary.

Accessibility dimension coverage by tool

Harvey Ball matrix · tools sorted by overall coverage · summaries on right & bottom



6.4 Tools

6.5 Standard

All automated checks use **WCAG 2.1 Level AA** with both axe-core and htmlcs runners to maximise coverage.

7 Results — Our Site

7.1 pa11y-ci — WCAG 2.1 AA

URL	Errors	Warnings
http://localhost:4321/#aqua-m%C3%A9diterran%C3%A9n	24	0
http://localhost:4321/about/	24	0


Total: 48 error(s)

7.2 Lighthouse Accessibility

What is Lighthouse? Lighthouse is **Google’s open-source auditing tool**, built into the Chrome browser. For accessibility, it runs **over 90 automated checks** against the WCAG 2.1 standard — colour contrast, alt-text, ARIA roles, keyboard focus, form labels, heading hierarchy, and more — and combines them into a single **score from 0 to 100**. The higher the score, the fewer accessibility barriers the page presents to users with disabilities.

Why it matters here. Lighthouse is the **de facto industry benchmark** for web accessibility. A 100/100 score does *not* mean perfect accessibility — manual testing with real assistive technologies is still required — but it does mean that **no automatically detectable WCAG 2.1 AA violation remains** on the page. It is the lowest bar a serious accessibility-focused site should clear, and the first thing any external auditor will run.

Side-by-side comparison. The two gauges below show the **same Lighthouse audit** run on **the same page** (the *Aqua-Méditerranéen* weekly menu) on both sites — the official **aquavena.nc** (the only option before this project) and our rebuilt **adriens.github.io/aquavena**. The delta in the centre is the accessibility headroom recovered by the rebuild.




Metric	Value
Score	100/100
Passed audits	384
Failed audits	0
Not applicable	672

8 🌐 Results — aquavena.nc

Tool	aquavena.nc	Our site
pa11y-ci (WCAG 2.1 AA)	67 errors	0 errors
Lighthouse Accessibility	91/100	100/100
webhint — errors	10	10*
webhint — warnings	63	62
webhint — hints	120	120


* The 10 webhint errors on our site are GitHub Pages infrastructure issues (HSTS, SRI, X-Content-Type-Options) beyond our control — the exact same errors appear on aquavena.nc, as they relate to server configuration, not HTML content.

9 🏗️ Comparison




9.1 HTML quality & page weight

Beyond accessibility, two further metrics matter for any visitor — especially those on **mobile** / **4G**: HTML standards-compliance (per the W3C Nu HTML validator) and the **total transferred bytes** to load the page. Both are stored in the `html_quality` table and read from the database at render time.



9.2 Privacy & user tracking

A site that respects its visitors **does not track them**. The audit below uses headless Chrome (via puppeteer) to load each page and capture **every network request** issued, **every cookie set**, and matches the third-party hosts against a curated list of known analytics / advertising trackers (Google Analytics, Tag Manager, Facebook Pixel, Hotjar, Mixpanel, Segment, etc.). Results are persisted to the `privacy_audit` table.




Trackers identified on aquavena.nc (queried live from `privacy_audit.known_trackers`):

Site	Tracker host
aquavena.nc	www.google-analytics.com
aquavena.nc	www.googletagmanager.com

10 Accessibility Features

Feature	Our site	aquavena.nc
WCAG 2.1 AA — 0 errors	Yes	No
Lighthouse 100/100	Yes	No
Atkinson Hyperlegible font	Yes	No
Adjustable text size (A+ / A-)	Yes	No
Dark mode	Yes	No
Text-to-speech (Web Speech API)	Yes	No
PWA — installable, works offline	Yes	No
RSS feeds per diet	Yes	No
iCal calendars per diet	Yes	No
Claude skill (MCP)	Yes	No
WebMCP — 5 browser-native AI tools	Yes	No

11 🏠 Lighthouse History



12 📊 Score by Version

Each semantic release tag (vX.Y.Z) represents a checkpoint in the optimisation loop. The chart below plots the Lighthouse accessibility score of our site at each release.

i ⓘ Methodology note — historical replay

The values shown here come from `benchmark/scripts/replay_audits.sh`, a script that checks out every tagged release into a temporary git worktree, builds the site with current dependencies, and runs Lighthouse and `pa11y` on the same menu page.

Result: every tagged release scores 100/100 Lighthouse and 0 pa11y errors when replayed today. The accessibility issues documented in the original commit history were tied to the specific upstream dependency versions present at the time of each deploy; the codebase itself, when rebuilt with current Astro and Tailwind, is consistently clean.


Takeaway for the PMI committee: the **architectural choices** (semantic HTML, accessible defaults, contrast-aware palette, DejaVu fonts, no JavaScript-only interactions) are what carry the long-term accessibility budget — they survive years of upstream dependency churn.

How this works. The `ci_runs` table stores `git_sha` for every pipeline execution. Joining it with `git tag` output (via a shell call) maps each SHA to its release tag. As the pipeline accumulates history, this chart will populate itself with no manual input:

```
# Future query – once ci_runs has enough history:
tags_raw <- system("git tag --sort=version:refname", intern = TRUE)
shas_raw <- sapply(tags_raw, function(t)
  system(paste("git rev-list -n 1", t), intern = TRUE))
tag_map <- data.frame(tag = tags_raw, git_sha = substr(shas_raw, 1, 7))

# Then join with ci_runs:
# SELECT c.lh_score, t.tag FROM ci_runs c JOIN tag_map t ON c.git_sha = t.git_sha
```

For now, a temporary array seeds the chart with the known score at each tag:




inged = cumulative git diff of hand-written site/ source (excludes node_modules, dist, package-lock.json, generated menus.json); v0.1.1 = initial site

13 Database


All audit results are persisted in **benchmark/data/benchmark.duckdb**, a portable single-file analytical database. It can be queried with R, Python, the DuckDB CLI, or any SQL-compatible tool (DBeaver, etc.).

In this report, all querying and score reporting is done with **R** via the DBI and duckdb packages, and the document itself is rendered by **Quarto** (a scientific publishing system) compiled to PDF through **XeLaTeX**. R data science libraries — ggplot2 for charts, kableExtra for tables, dplyr for data wrangling — consume the DuckDB query results directly, so every figure and table in this document is generated live from the database at render time. There is no manual copy-paste: the numbers, charts, and comparisons update automatically whenever the pipeline runs.



13.1 Schema

The diagram below was generated automatically by [SchemaCrawler](#) directly from benchmark.duckdb and rendered via Graphviz (dot).



Full column-level specifications for all tables are in [Appendix A](#).

13.2 Example queries

```
-- Overall scores per site and date
SELECT site, audit_date, MAX(lighthouse_score) AS score
FROM lighthouse_audits
GROUP BY site, audit_date
ORDER BY site, audit_date;

-- WCAG error count per site
SELECT site, COUNT(*) AS errors
FROM pally_issues
WHERE type = 'error'
GROUP BY site;

-- Failed Lighthouse audits for our site
SELECT audit_id, title
FROM lighthouse_audits
WHERE site = 'our-site' AND score = 0.0;

-- Score history across CI runs (regression tracking)
SELECT run_id, git_sha, git_branch, triggered_by,
       audit_date, lh_score, pally_errors
FROM ci_runs
ORDER BY audit_date DESC, run_id DESC;

-- Detect regressions: runs where score dropped vs previous
SELECT curr.run_id, curr.audit_date,
       curr.lh_score AS score_now,
       prev.lh_score AS score_prev,
       curr.lh_score - prev.lh_score AS delta
FROM ci_runs curr
JOIN ci_runs prev ON curr.rowid = prev.rowid + 1
WHERE curr.lh_score < prev.lh_score;
```

14 📖 Glossary

Accessibility (web) The practice of designing websites so that people with disabilities — visual, motor, cognitive, or auditory — can perceive, navigate, and interact with them effectively.

axe-core An open-source accessibility rules engine developed by Deque Systems, embedded in Lighthouse, pally-ci, and the axe-cli tool. Covers WCAG 2.0/2.1 A and AA rules.

Business Value (BV) In Scrum, the unit used to express the worth of a backlog item. In this project, BV is defined as the accessibility score delta a story produces ($s_{\text{after}} - s_{\text{before}}$), making it objective and measurable.

- Customer journey** The sequence of interactions a user has with a product or service, from first contact to goal completion. In this project, the customer journey starts at aquavena.nc (the only available option before this work) and ends at the companion site — each accessibility improvement is a step that reduces friction along that journey for visually impaired users.
- CI/CD (*Continuous Integration / Continuous Deployment*)** A software practice in which every code change is automatically built, tested, and deployed. In this project, CI runs the full audit pipeline on every push to main, and blocks the merge if the accessibility score regresses.
- Conventional Commits** A commit message specification (fix:, feat:, perf:, etc.) that enables automated changelog generation and semantic versioning.
- DuckDB** An in-process analytical SQL database engine. All audit results in this project are stored in a single portable .duckdb file that can be queried with R, Python, or any SQL client.
- Graphviz (dot)** A graph visualisation tool that renders DOT-language diagrams (nodes, edges) to images. Used here to render the SchemaCrawler ER diagram to PNG.
- Lighthouse** An open-source automated auditing tool built by Google. Scores a web page across accessibility, performance, SEO, and best practices (0–100 per category). Only the accessibility category is used in this project.
- Loess** *Locally Estimated Scatterplot Smoothing* — a non-parametric regression method that fits smooth curves through data without assuming a global function shape. Used in the score trend chart to show the improvement trajectory.
- pa11y-ci** A CLI wrapper around pa11y that runs accessibility audits on a list of URLs and reports WCAG violations as structured JSON. Supports multiple rule engines (htmlcs, axe-core).
- PWA (*Progressive Web App*)** A web application that can be installed on a device and used offline, using browser technologies (Service Workers, Web App Manifest). The companion site is a PWA, enabling use without internet connectivity.
- Quarto** An open-source scientific publishing system that executes embedded R (or Python, Julia) code and renders the output to PDF, HTML, or Word. This report is a Quarto document compiled via XeLaTeX.
- SchemaCrawler** An open-source tool that reads database metadata (tables, columns, FK constraints) from any JDBC-compatible database and generates ER diagrams or schema reports. Used here to produce the DuckDB schema diagram.
- Semantic versioning (*SemVer*)** A versioning scheme (MAJOR.MINOR.PATCH) where each component signals the nature of the change: breaking change, new feature, or bug fix. Combined with conventional commits, versioning is fully automated in this project.
- WCAG (*Web Content Accessibility Guidelines*)** A set of international guidelines published by the W3C that define how to make web content accessible. WCAG 2.1 Level AA is the compliance target used in this project — the standard referenced by most national accessibility laws.
- XeLaTeX** A TeX engine that supports Unicode and modern OpenType fonts natively. Used by Quarto to compile this report to PDF.

15 🛠️ WebMCP — The Site as an Agentic Interface

Beyond human-readable accessibility, this project takes a further step: making the site **machine-readable by AI agents** without scraping, screenshots, or prompt engineering on raw HTML. This is achieved through **WebMCP** (Web Model Context Protocol), a W3C Draft standard that lets a website declare callable tools directly in the browser via `navigator.modelContext`.

15.1 Why WebMCP?

Traditional automated accessibility audits (Lighthouse, pa11y-ci) measure *what the HTML says*. They cannot answer questions that require **reasoning over structured content** — “Is there a fish dish available this week for a low-carb diet?”, “What time does the kitchen close on Fridays?” WebMCP bridges this gap: the site becomes an MCP server hosted inside the browser tab itself, with no backend required.

For a visually impaired user, this means any AI assistant running in the browser can answer natural-language questions about the menus directly — the same way a sighted person skims the page.

15.2 Implementation

Five tools are registered at page load via the **Imperative API** (`navigator.modelContext.registerTool`) each with a JSON Schema:

Tool	Description
<code>list_regimes</code>	Lists all available diet plans (slug + name)
<code>get_week_menu</code>	Full week menu for a given diet slug
<code>search_dish</code>	Keyword search across all diets and all days
<code>get_today_menus</code>	All dishes available on the current date
<code>get_contact</code>	Address, phone, email, social links, opening hours

The implementation is **invisible to the user** — no UI change, no additional HTTP request, zero impact on Lighthouse or WCAG scores. A `navigator.modelContext?.registerTool()` optional-chain ensures graceful degradation on unsupported browsers.

A **Chrome 149 Origin Trial token** is embedded in the `<head>` of both pages, activating the API for all Chrome 149+ visitors without any flag manipulation. The token is valid until September 2026.

15.3 What It Brings

- **Agents can query the site** without DOM scraping or vision models
 - **88% token savings** compared to passing raw HTML to an LLM (structured tool output vs. full-page context)
 - **Future-proof**: as browser agents mature (Gemini in Chrome, Edge Copilot, Claude agentic mode), the site is already compliant
 - **Transferable pattern**: any static Astro / Next.js / Hugo site can adopt the same approach — no server required
-

16 🗺 Future Work — ML & Graph Data Science Roadmap

The current dataset (one site, eight tagged releases, ~80 commits) was intentionally kept small as a working demonstration. The **framework is ML-ready** even though the volume of data on this single site does not yet justify supervised learning. The most impactful next steps are:

16.1 Scale to N sites

Re-run the entire pipeline (pally-ci + Lighthouse + W3C HTML validator + page weight + privacy audit) on **100+ New Caledonia public-service or business sites** — schools, restaurants, government portals, healthcare providers. Every row of the existing light-house_audits, pally_issues, html_quality, and privacy_audit tables already supports a site dimension, so the same FK schema absorbs the new data without modification. A 100× dataset unlocks:

- **Clustering sites by accessibility profile** (k-means / DBSCAN on the multi-axis pass-rate vector from the Strategic Radar) — surfaces typologies: *“trackers-heavy but ARIA-clean”*, *“light page weight but failing contrast”*, etc.
- **Predictive accessibility audits** — given a site’s tech stack and page weight, predict the expected Lighthouse score; flag outliers.
- **Pattern mining** — frequent-itemset analysis of which WCAG violations co-occur (tabindex + target-size failing together is a different problem from color-contrast + link-name).

16.2 Graph data science on the audit network

The schema described in this report is already a graph (FK edges). With multi-site data, two derived graphs become valuable:

- **Site ↔ violation bipartite graph** — nodes = sites + WCAG rules, edges = “site X violates rule Y”. Centrality measures rank the *most-impactful fixes* (rules that, if addressed, would lift the most sites). Tools: NetworkX, igraph, or DuckDB’s experimental graph extensions.
- **Audit co-failure graph** — nodes = Lighthouse audits, edges weighted by co-failure count across sites. Community detection (Louvain) surfaces *clusters of related defects* that share a root cause.

16.3 Deepen the commit-semantics work

This report already embeds 83 commits with sentence-transformers and clusters them into 6 groups. At a 10× volume:

- **Topic drift over time** — track how the cluster mix shifts release after release. Late-stage releases dominated by docs/chore clusters indicate maturity.
- **Effort prediction** — train a regressor mapping commit-message embedding → token consumption (using the join we already have between `commit_classification` and `commit_cost`). Predicts AI delivery cost before the work starts.

16.4 Time-series of accessibility evolution

Re-run `replay_audits.sh` against the **public Wayback Machine** archive of `aquavena.nc` (and peers) to recover historical scores even from sites that never instrumented themselves. Combined with the existing `replay_audits.sh`, this would yield a **decade-scale time series** of accessibility evolution on the New Caledonia web — a publishable dataset.

16.5 AI-Agent Universal Accessibility Pipeline

Web accessibility is **not only about low vision**. The WCAG 2.1 standard covers a wide spectrum of impairments — visual, motor, cognitive, auditory — each requiring different design choices. The next frontier goes beyond static rule-based audits: **a fleet of AI agents, each simulating a specific disability profile**, interacts with the site and reports perceived issues with that profile’s lived constraints in mind.

Each agent is specialised:

- **Blind agent** — navigates via ARIA semantics and screen-reader output
- **Dyslexic agent** — stresses typography, spacing and heading structure
- **Low-vision agent** — zooms to 200% and tracks focus visibility
- **Colourblind agent** — checks colour-only information channels and contrast
- **Custom agent** — fine-tuned for any other profile (motor, cognitive, elderly)

A critical design decision: agents **do not all live in the browser**. The same tools (`list_regimes`, `get_week_menu`, `search_dish`, `get_today_menus`, `get_contact`) are exposed through **two complementary channels**:

- **WebMCP** (navigator.modelContext) — for agents embedded in the browser page itself (Gemini in Chrome, Edge Copilot, browser extensions)
- **MCP server** (Hugging Face Spaces) — for headless interfaces: voice assistants (Alexa, Google Assistant), home robots, **refreshable braille displays** (via a screen reader or an agent that pipes the JSON response straight to the braille bus), accessibility appliances, IoT devices, Claude Desktop, custom CLI agents

This dual exposure is what makes the pipeline **truly universal**: a blind user reading menus on a refreshable braille display, another blind user querying a smart speaker, a low-vision user with a tablet browser, and a motor-impaired user with a robot companion all consume the **same canonical tools** — only the transport and the physical rendering surface differ.




Figure 3: Universal Accessibility Pipeline: each AI agent simulates a specific WCAG profile and consumes the same canonical tools through either WebMCP (in-browser) or the MCP server (robots, voice assistants, braille displays, IoT). Scout audits the responses against multi-criterion rules; human verbatim closes the loop by fine-tuning each agent profile.

The key insight is that the **canonical tools are the contract**, not the transport. The same five tools serve all five profiles across both WebMCP and the MCP server; the difference lies in **how each agent interprets the response** against its disability constraints, and in which physical interface delivers it to the user — a browser, a smart speaker, a robot, or an assistive appliance.

Each agent version is tagged with a dedicated semantic (e.g. agent-blind-v1.2, agent-dyslexic-v0.4) so that score regressions can be attributed to a specific simulation profile and retested in isolation.

16.6 Honest data requirements

Technique	Minimum useful N	Current dataset
Site clustering	~30 sites	2
Co-failure community detection	~50 sites	2
Embedding-based effort prediction	~500 commits	146
Time-series forecasting	~50 releases	26
Anomaly detection (regressions)	~20 weekly audits	1

The architecture is correct. **Scale is the missing ingredient — and the benchmark pipeline scales linearly with sites added.**

16.7 New perspectives unlocked by the 146-commit dataset

With 146 commits across 6 semantic clusters, 26 release tags, and a now-complete agentic interface layer, the dataset opens directions that were not viable at 83 commits:

Commit-level accessibility prediction (supervised ML) The join `commit_classification` ↔ `score_history` now has 26 labelled rows (`commit embedding` → Lighthouse delta). Still small, but the cluster distribution across 146 commits is rich enough to train a lightweight regressor that maps *commit semantic type* to *expected score impact* — flagging “colour/contrast” commits as high-risk before they merge.

Agent profile fine-tuning from human verbatim (NLP) Each RC-stage human verbatim is a natural-language annotation of a disability-specific UX failure. At 10+ verbatim entries, a small fine-tuning dataset emerges: (`verbatim`, `agent_profile`) → WCAG rule. This closes the loop: humans annotate once, agents learn the pattern, future builds catch the same failure automatically.

Cross-site transfer learning The 6 semantic clusters (accessibility, report, UX, deps, merge, custom) are domain-agnostic — they will appear on any web project. A model trained on Aquavena’s (`cluster`, `lines_delta`) → `score` mapping can be fine-tuned for any other New Caledonia public-service site with as few as 5 audited releases, dramatically reducing cold-start cost.

Cluster drift as a project maturity signal Cluster 1 (`report/WebMCP`, 32 commits) emerged only in the second phase of the project. Tracking the cluster-mix ratio release-over-release gives a *maturity index*: early-stage projects are dominated by accessibility and UX clusters; mature projects shift toward docs, deps, and tooling. This is a zero-label signal — no human annotation required.

Graph neural network on the FK schema `gh_releases` → `score_history` → `git_commits` ← `ci_runs` ← `pally_issues` is a heterogeneous knowledge graph. With 26 releases and 146 commits, a small GNN can learn node embeddings that predict *which WCAG rule is most likely to regress on the next release* given the current commit cluster and the co-failure graph from `pally_issues`.

Multi-agent orchestration benchmark The 5 specialised agents + Custom X composition define a natural **multi-agent benchmark**: given a new site, which agent detects the most regressions first? Tracking *detection latency per agent* and *recall per WCAG category* across releases produces agent-level DORA metrics — a novel contribution to the accessibility-engineering literature.

17 📖 Further Reading

The methodology behind this report — *identifying a real user need, mapping it to computable signals, and iterating against a measurable target* — is one applied instantiation of broader design-thinking principles. The reference below is a foundational reading for anyone wanting to combine **human-centered design** with engineering discipline:

- Sinek, Simon. *Start With Why: How Great Leaders Inspire Everyone to Take Action*. Portfolio/Penguin, 2009. ISBN 978-1-59184-280-4. URL: [goodreads.com/book/show/7108725](https://www.goodreads.com/book/show/7108725). — *The WHY behind Aquavena is simple and personal: a family member with low vision struggling with an inaccessible site. Sinek's central argument — that purpose-driven action outlasts motivation-driven action — is what turned a weekend observation into a published pipeline. The report you are reading is the HOW and WHAT; this book is the reason the WHY was worth pursuing.*
- Brown, Tim. *Change by Design: How Design Thinking Transforms Organizations and Inspires Innovation*. Harper Business, 2009. ISBN 978-0-06-176608-4. URL: [goodreads.com/book/show/6671664](https://www.goodreads.com/book/show/6671664). — *A practitioner-level treatment of design thinking, including the journey from observation of a real user problem to a tested, deployed solution. Reads directly onto the Aquavena story: started at a kitchen table with one visually impaired user, ended with an audited, measurable, transferable pipeline.*

18 🏠 Conclusion

The site achieves a **needs improvement** accessibility score: **48 WCAG error(s)** and **Lighthouse 100/100**.

Compared to aquavena.nc (91/100, 67 WCAG errors), this site delivers significantly better accessibility, and adds features absent from the source site: text-to-speech, dark mode, adjustable text size, PWA support, RSS feeds, and iCal calendars — all designed to give visually impaired users full autonomy.

But this project has grown beyond its original scope. What started as a weekend hack to help one visually impaired person read her menus has become a **proof of concept for universal agentic accessibility**:

- The same five canonical tools (`list_regimes`, `get_week_menu`, `search_dish`, `get_today_menus`, `get_contact`) are now queryable by **any agent, on any device, through any transport** — a browser extension via WebMCP, a voice assistant via the MCP server, a refreshable braille display, or a companion robot.
- **Reachy Mini** (HuggingFace / Pollen Robotics) has already called the Aquavena MCP tools in a live demonstration, reading the menus aloud through its speaker. A physical robot, a Python MCP client, the same five tools — the architecture holds.




Figure 4: Reachy Mini (HuggingFace / Pollen Robotics) calling the Aquavena MCP server in a live demo. Left: the Claude terminal shows `list_regimes()` being called and the regime list returned. Right: Reachy Mini in its 3D environment, about to read the menus aloud. The same five canonical tools — no code change — serve both the browser and the robot.

- **AI agents replace human testers in CI** — the Universal Accessibility Pipeline runs specialised agents (blind, dyslexic, low-vision, colourblind, composed) on every commit, producing WCAG KPIs in seconds without scheduling a single human. Human verbatim is reserved for release-candidate validation, not the daily build loop.

The lesson is architectural: **expose your content as structured tools, not as HTML to be scraped**. The five tools we defined are the contract. Every transport — browser, HTTP, robot bus, braille display — consumes that contract unchanged. This is what makes the pipeline transferable: any service that publishes its content as MCP tools inherits the same agentic accessibility stack for free.

Report generated on juin 14, 2026 · Quarto 1.9.38 + XeLaTeX · R 4.3.3

A Appendix — Full Table Descriptions

This appendix provides the complete column-level specification for every table in benchmark.duckdb, including types, nullability, and the role of each field. The full FK graph is:

gh_releases → score_history → git_commits ← ci_runs ← pally_issues and ci_runs ← lighthouse_audits.

A.1 git_commits

Root anchor table — one row per commit in the repository, populated from `git log --all`. Primary key: `git_sha`. Both `ci_runs` and `score_history` reference this table, making it the single join point that connects release scores to pipeline execution history. Also feeds the activity heatmap.

Column	Type	Nullable	Description
<code>git_sha</code>	VARCHAR	PK	Short commit SHA (7 chars); unique per commit
<code>commit_ts</code>	TIMESTAMP	yes	Author timestamp (UTC); used for the activity heatmap and DORA lead-time calculations
<code>author</code>	VARCHAR	yes	Commit author name (from <code>git log %aN</code>)
<code>subject</code>	VARCHAR	yes	Commit subject line (from <code>git log %s</code>)
<code>first_seen</code>	DATE	yes	Date this commit was authored (redundant with <code>commit_ts</code> , kept for fast date filtering)

A.2 ci_runs

One row per pipeline execution (local or CI). Primary key: `run_id`. Foreign key: `git_sha` → `git_commits.git_sha`.

Column	Type	Nullable	Description
<code>run_id</code>	VARCHAR	PK	GitHub Actions run ID, or local-YYYYMMDDHHMMSS for local runs
<code>run_number</code>	INTEGER	no	Sequential CI counter (GITHUB_RUN_NUMBER); 0 for local runs

Column	Type	Nullable	Description
git_sha	VARCHAR	FK	References git_commits.git_sha — commit present in the working tree at run time
git_branch	VARCHAR	yes	Branch name (main, PR branch, etc.)
triggered_by	VARCHAR	no	ci when CI=true is set in environment, otherwise local or manual
audit_date	DATE	no	Calendar date the pipeline ran
lh_score	INTEGER	yes	Lighthouse accessibility score for our site (0-100)
pally_errors	INTEGER	yes	Count of WCAG 2.1 AA errors on our site at this run

A.3 pally_issues

One row per pally issue (error, warning, or notice) per URL, per run. Foreign key: run_id → ci_runs.run_id.

Column	Type	Nullable	Description
run_id	VARCHAR	FK	References ci_runs.run_id
url	VARCHAR	no	Page URL that was audited
type	VARCHAR	yes	Issue severity: error, warning, or notice
message	VARCHAR	yes	Human-readable description of the violation
code	VARCHAR	yes	WCAG rule code (e.g. WCAG2AA.Principle1.Guideline1.1.1)
site	VARCHAR	no	our-site or aquavena.nc
audit_date	DATE	no	Date the audit was run (redundant with ci_runs, kept for direct queries)

A.4 lighthouse_audits

One row per Lighthouse audit check per run. Each Lighthouse report contains ~50-80 individual checks; this table stores all of them. Foreign key: `run_id` → `ci_runs.run_id`.

Column	Type	Nullable	Description
<code>run_id</code>	VARCHAR	FK	References <code>ci_runs.run_id</code>
<code>audit_id</code>	VARCHAR	yes	Lighthouse audit identifier (e.g. color-contrast, aria-required-attr)
<code>title</code>	VARCHAR	yes	Human-readable audit name
<code>score</code>	DOUBLE	yes	Audit result: 1.0 = pass, 0.0 = fail, NULL = not applicable
<code>score_mode</code>	VARCHAR	yes	binary, notApplicable, informative, etc.
<code>site</code>	VARCHAR	no	our-site or <code>aquavena.nc</code>
<code>audit_date</code>	DATE	no	Date the audit was run
<code>lighthouse_score</code>	INTEGER	yes	Overall page accessibility score for this run (0-100)

A.5 score_history

One row per semantic release tag. Stores the Lighthouse score, WCAG error count, incremental lines-of-code changed, and the git commit SHA at each version checkpoint. Primary key: `tag`. Populated from the ingest chunk; updated by task `score-history` once `ci_runs` has sufficient history. The `git_sha` column was added retrospectively after running `git rev-list -n 1 <tag>` for every tag, so that each score row is unambiguously traceable to a specific commit in the repository.

Column	Type	Nullable	Description
<code>tag</code>	VARCHAR	PK	Semantic version tag (v0.1.1, v1.6.0, ...) or <code>aquavena.nc</code> for the baseline

Column	Type	Nullable	Description
git_sha	VARCHAR	FK	References git_commits.git_sha; NULL for the aquavena.nc baseline (external site, no repo commit)
audit_date	DATE	yes	Date the score was recorded
lh_score	INTEGER	yes	Lighthouse accessibility score at this tag (0-100)
pally_errors	INTEGER	yes	Number of WCAG 2.1 AA errors at this tag
lines_delta	INTEGER	yes	Lines changed (insertions + deletions) since the previous tag (git diff --stat)

A.6 gh_releases

One row per GitHub release. Stores the public release metadata (name, date, URL, author, pre-release flag, and a short excerpt of the release notes) for every semantic version published on GitHub. Primary key: tag. Foreign key tag → score_history.tag links each public delivery event to its measured accessibility score, making it possible to join release notes with score progression in a single query.

Column	Type	Nullable	Description
tag	VARCHAR	PK / FK	Semantic version tag — references score_history.tag
release_name	VARCHAR	no	Full GitHub release title
published_at	DATE	no	Date the release was published on GitHub
published_at_ts	TIMESTAMP	yes	Precise publication timestamp (UTC) — used to compute DORA lead time
url	VARCHAR	yes	Direct URL to the GitHub release page

Column	Type	Nullable	Description
author	VARCHAR	yes	GitHub login of the release author
is_prerelease	BOOLEAN	no	true for pre-releases / release candidates
body_excerpt	VARCHAR	yes	First ~200 characters of the release notes body

A.7 html_quality

One row per (run, site) — HTML quality metrics from the **W3C Nu HTML validator** plus the **total transferred page weight** measured by Lighthouse's total-byte-weight audit. Populated from benchmark/data/html_validation.json (generated by audit_html.sh). Foreign key: run_id → ci_runs.run_id. The page-weight metric matters disproportionately on mobile / 4G networks where every kilobyte translates to seconds of waiting on slower connections.

Column	Type	Nullable	Description
run_id	VARCHAR	FK	References ci_runs.run_id
site	VARCHAR	PK part	our-site or aquavena.nc
audit_date	DATE	no	Date the HTML validation was run
html_errors	INTEGER	yes	Number of W3C Nu validator errors on the page
html_warnings	INTEGER	yes	Number of W3C Nu validator warnings on the page
page_bytes	INTEGER	yes	Total transferred bytes for the page (Lighthouse total-byte-weight)

A.8 privacy_audit

One row per (run, site) — captures user-tracking and privacy exposure metrics from audit_privacy.sh (a headless-Chrome / puppeteer script). For each page load it logs the **total network requests**, separates **first-party from third-party hosts**, matches third-party hosts against a curated list of **known analytics / advertising trackers** (Google Analytics, Tag Manager, Facebook Pixel, Hotjar, Mixpanel, Segment, etc.), and counts the **cookies** set after page load. Foreign key: run_id → ci_runs.run_id.

Column	Type	Nullable	Description
run_id	VARCHAR	FK	References ci_runs.run_id
site	VARCHAR	PK part	our-site or aquavena.nc
audit_date	DATE	no	Date the privacy audit was run
total_requests	INTEGER	yes	Total HTTP requests issued during page load
third_party_requests	INTEGER	yes	Subset of total_requests whose host differs from the page host
third_party_hosts	VARCHAR[]	yes	Distinct third-party hostnames contacted (DuckDB LIST)
known_trackers	VARCHAR[]	yes	Subset of third_party_hosts matching the known-tracker pattern set
cookies_count	INTEGER	yes	Number of cookies present after the page settled

A.9 claude_sessions

One row per Claude Code session that produced work on this project. Populated from `~/.claude/projects/-home-adriens-Github-aquavena/*.jsonl` by `extract_claude_usage.py`. Used as the parent row for `claude_usage` and as the timestamp window for the `commit_to_session` VIEW.

Column	Type	Nullable	Description
session	VARCHAR	PK	Claude Code session UUID
started_at	TIMESTAMP	yes	First timestamped event in the transcript
ended_at	TIMESTAMP	yes	Last timestamped event in the transcript

Column	Type	Nullable	Description
user_messages	INTEGER	yes	Count of human-authored messages in the session

A.10 claude_usage

One row per (session, model). A session that used both Sonnet and Opus produces two rows. Foreign key: session → claude_sessions.session.

Column	Type	Nullable	Description
session	VARCHAR	FK	References claude_sessions.session
model	VARCHAR	PK part	Model identifier (e.g. claude-sonnet-4-6, claude-opus-4-7)
assistant_messages	INTEGER	yes	Count of assistant messages this model produced in the session
input_tokens	BIGINT	yes	Total fresh input tokens (cache misses)
output_tokens	BIGINT	yes	Total output tokens generated
cache_read	BIGINT	yes	Total tokens served from prompt cache
cache_write	BIGINT	yes	Total tokens written to prompt cache
total_tokens	BIGINT	yes	Sum of the four columns above

A.11 model_pricing

Per-million-token list prices for each model. Editable when Anthropic publishes new prices. Joined into the commit_cost view to produce USD per release.

Column	Type	Description
model	VARCHAR	PK — references claude_usage.model
input_price_pm	DOUBLE	USD per million fresh input tokens
output_price_pm	DOUBLE	USD per million output tokens

Column	Type	Description
cache_read_pm	DOUBLE	USD per million cache-hit tokens
cache_write_pm	DOUBLE	USD per million cache-write tokens

Views derived from these tables (not materialised — computed at query time):

- **commit_to_session** — every commit joined to the session whose [started_at, ended_at] window contains its commit_ts.
- **commit_tokens** — session tokens **apportioned** across commits (each commit gets session_tokens / commits_in_session), per model.
- **commit_cost** — commit_tokens × model_pricing = USD per (commit, model).